

CS6240

Rachna Reddy Melacheruvu

HW 4

<https://github.ccs.neu.edu/melacheruvur/cs6240/tree/master/Assignment4/PR-Spark/Spark-Demo>

<https://github.ccs.neu.edu/melacheruvur/cs6240/tree/master/Assignment4/PR-MapReduce/MR-Demo>

Citation:

1. Changing of jobs – <https://coe4bd.github.io/HadoopHowTo/multipleJobsSingle/multipleJobsSingle.html>
2. NLineInputFormat - <https://stackoverflow.com/questions/28871899/when-to-use-nlineinputformat-in-hadoop-map-reduce>
3. To setValue of counter from string to long - <https://stackoverflow.com/questions/34403552/hadoop-counters-how-do-i-use-different-type-of-counters>
4. To solve the error of intermediate output files not being able to access on s3 - <https://github.com/51zero/eel-sdk/issues/356>

URLs:

Page Rank Spark-

1. **Project:**
<https://github.ccs.neu.edu/melacheruvur/cs6240/tree/master/Assignment4/PR-Spark/Spark-Demo>
2. **Logs:** <https://github.ccs.neu.edu/melacheruvur/cs6240/blob/master/Assignment4/PR-Spark/Spark-Demo/logs>
3. **Output:**
<https://github.ccs.neu.edu/melacheruvur/cs6240/tree/master/Assignment4/PR-Spark/Spark-Demo/output>

Page Rank MapReduce-

1. **Project:**
<https://github.ccs.neu.edu/melacheruvur/cs6240/tree/master/Assignment4/PR-MapReduce/MR-Demo>
2. **Logs :**
 1. 6 machines – <https://github.ccs.neu.edu/melacheruvur/cs6240/tree/master/Assignment4/PR-MapReduce/MR-Demo/1000kfor6/logs>
 2. 11 machines - <https://github.ccs.neu.edu/melacheruvur/cs6240/blob/master/Assignment4/PR-MapReduce/MR-Demo/1000kfor11/Logs/syslog.txt>
3. **Output:**

1. **6 machines** – <https://github.ccs.neu.edu/melacheruvur/cs6240/blob/master/Assignment4/PR-MapReduce/MR-Demo/1000kfor6/output/outputFinal/part-r-00003>
2. **11 machines** – <https://github.ccs.neu.edu/melacheruvur/cs6240/blob/master/Assignment4/PR-MapReduce/MR-Demo/1000kfor11/Output/outputFinal/part-r-00016>

Page Rank in Spark

Pseudo-code

```
val k: Int = args(0).toInt // k as input
val iters: Int = 10 // number of iterations to be performed
val maxK: Int = k*k // k^2 required to calculate initial Page Rank
val initialRank: Float = (1/maxK.toFloat)
val vertices = List.range(1,maxK+1) // list of vertices from 1 to k
val graph = sc.parallelize(vertices.map(vertex =>{ // creates graph RDD as 1-> 2 and for every k th vertex
// it points to dummy vertex as k -> 0

  if(vertex%k==0)
    (vertex,0)
  else
    (vertex,vertex+1)
}))
val dummyRDD = List(0).map(x=>(x,0.0f))
val verRanks = vertices.map(vertex=>(vertex,initialRank)) // to calculate ranks for each vertex with initial PR
val ranks = sc.parallelize(verRanks.union(dummyRDD))

for (i <- 1 to iters) { // Page Rank Algorithm Iteration

  val contribs = graph.join(ranks) // Join graph and ranks so we get vertex, adjacency list
  .values.map{ // and page rank
    case(u,v)=>(u,v) // store only vertex and it's incoming PR
  }
  val combineRanks = contribs.reduceByKey(_ + _) // add all incoming PR for each vertex
  val dummyRank = combineRanks.lookup(0).head.toFloat // handling PageRank Mass by distributing it to all vertices
  val PRMass = dummyRank/maxK
  val nextRanks = ranks.leftOuterJoin(combineRanks).map({ // to add new PR to each vertex and if vertex has no
    case (k,(u,None)) => // incoming link, then just add the dangling mass
    {
      if(k==0)
        (k,0.0f)
      else
        (k,PRMass)
    }
    case (k,(u,Some(v))) =>
    {
      if(k==0)
        (k,0.0f)
      else
        (k,v+PRMass)
    }
  })
  ranks = nextRanks // newly computed ranks are given to ranks variable for
```

```

// next iteration
var pagerankvalue = ranks.values.sum()
// calculating total PR of the graph
logger.info("Page Rank Value at-"+i+"th iteration is "+pagerankvalue)
}
logger.info(ranks.toDebugString)
val output = ranks.filter(_._1<=100)
// writing output to file of 100 vertices
output.saveAsTextFile(args(1))

```

Actions in Program

```

saveAsTextFile
sum
lookup

```

Ran program for k=100 and 10 iterations. Output shows PageRanks for vertices 0 to 100. Output in above provided URL

Lineage for RDD Ranks

1. For 1 iteration

```

(4) MapPartitionsRDD[11] at map at PageRank.scala:59 []
| MapPartitionsRDD[10] at leftOuterJoin at PageRank.scala:59 []
| MapPartitionsRDD[9] at leftOuterJoin at PageRank.scala:59 []
| CoGroupedRDD[8] at leftOuterJoin at PageRank.scala:59 []
+-(4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []
| ShuffledRDD[7] at reduceByKey at PageRank.scala:54 []
+-(4) MapPartitionsRDD[6] at map at PageRank.scala:49 []
| MapPartitionsRDD[5] at values at PageRank.scala:49 []
| MapPartitionsRDD[4] at join at PageRank.scala:48 []
| MapPartitionsRDD[3] at join at PageRank.scala:48 []
| CoGroupedRDD[2] at join at PageRank.scala:48 []
+-(4) ParallelCollectionRDD[0] at parallelize at PageRank.scala:33 []
+-(4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []

```

2. For 2 iterations

```

(4) MapPartitionsRDD[23] at map at PageRank.scala:59 []
| MapPartitionsRDD[22] at leftOuterJoin at PageRank.scala:59 []
| MapPartitionsRDD[21] at leftOuterJoin at PageRank.scala:59 []
| CoGroupedRDD[20] at leftOuterJoin at PageRank.scala:59 []
+-(4) MapPartitionsRDD[11] at map at PageRank.scala:59 []
| | MapPartitionsRDD[10] at leftOuterJoin at PageRank.scala:59 []
| | MapPartitionsRDD[9] at leftOuterJoin at PageRank.scala:59 []
| | CoGroupedRDD[8] at leftOuterJoin at PageRank.scala:59 []
| +--(4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []
| | ShuffledRDD[7] at reduceByKey at PageRank.scala:54 []
| +--(4) MapPartitionsRDD[6] at map at PageRank.scala:49 []
| | MapPartitionsRDD[5] at values at PageRank.scala:49 []
| | MapPartitionsRDD[4] at join at PageRank.scala:48 []
| | MapPartitionsRDD[3] at join at PageRank.scala:48 []
| | CoGroupedRDD[2] at join at PageRank.scala:48 []
| +--(4) ParallelCollectionRDD[0] at parallelize at PageRank.scala:33 []
| +--(4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []
| ShuffledRDD[19] at reduceByKey at PageRank.scala:54 []
+-(4) MapPartitionsRDD[18] at map at PageRank.scala:49 []
| MapPartitionsRDD[17] at values at PageRank.scala:49 []

```

```

| MapPartitionsRDD[16] at join at PageRank.scala:48 []
| MapPartitionsRDD[15] at join at PageRank.scala:48 []
| CoGroupedRDD[14] at join at PageRank.scala:48 []
+-(4) ParallelCollectionRDD[0] at parallelize at PageRank.scala:33 []
+-(4) MapPartitionsRDD[11] at map at PageRank.scala:59 []
| MapPartitionsRDD[10] at leftOuterJoin at PageRank.scala:59 []
| MapPartitionsRDD[9] at leftOuterJoin at PageRank.scala:59 []
| CoGroupedRDD[8] at leftOuterJoin at PageRank.scala:59 []
+-(4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []
| ShuffledRDD[7] at reduceByKey at PageRank.scala:54 []
+-(4) MapPartitionsRDD[6] at map at PageRank.scala:49 []
| MapPartitionsRDD[5] at values at PageRank.scala:49 []
| MapPartitionsRDD[4] at join at PageRank.scala:48 []
| MapPartitionsRDD[3] at join at PageRank.scala:48 []
| CoGroupedRDD[2] at join at PageRank.scala:48 []
+-(4) ParallelCollectionRDD[0] at parallelize at PageRank.scala:33 []
+-(4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []

```

3. For 3 iterations

```

(4) MapPartitionsRDD[35] at map at PageRank.scala:59 []
| MapPartitionsRDD[34] at leftOuterJoin at PageRank.scala:59 []
| MapPartitionsRDD[33] at leftOuterJoin at PageRank.scala:59 []
| CoGroupedRDD[32] at leftOuterJoin at PageRank.scala:59 []
+-(4) MapPartitionsRDD[23] at map at PageRank.scala:59 []
| MapPartitionsRDD[22] at leftOuterJoin at PageRank.scala:59 []
| MapPartitionsRDD[21] at leftOuterJoin at PageRank.scala:59 []
| CoGroupedRDD[20] at leftOuterJoin at PageRank.scala:59 []
+-(4) MapPartitionsRDD[11] at map at PageRank.scala:59 []
| MapPartitionsRDD[10] at leftOuterJoin at PageRank.scala:59 []
| MapPartitionsRDD[9] at leftOuterJoin at PageRank.scala:59 []
| CoGroupedRDD[8] at leftOuterJoin at PageRank.scala:59 []
| +-(4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []
| | ShuffledRDD[7] at reduceByKey at PageRank.scala:54 []
| +-(4) MapPartitionsRDD[6] at map at PageRank.scala:49 []
| | MapPartitionsRDD[5] at values at PageRank.scala:49 []
| | MapPartitionsRDD[4] at join at PageRank.scala:48 []
| | MapPartitionsRDD[3] at join at PageRank.scala:48 []
| | CoGroupedRDD[2] at join at PageRank.scala:48 []
| | +-(4) ParallelCollectionRDD[0] at parallelize at PageRank.scala:33 []
| | +-(4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []
| ShuffledRDD[19] at reduceByKey at PageRank.scala:54 []
+-(4) MapPartitionsRDD[18] at map at PageRank.scala:49 []
| MapPartitionsRDD[17] at values at PageRank.scala:49 []
| MapPartitionsRDD[16] at join at PageRank.scala:48 []
| MapPartitionsRDD[15] at join at PageRank.scala:48 []
| CoGroupedRDD[14] at join at PageRank.scala:48 []
+-(4) ParallelCollectionRDD[0] at parallelize at PageRank.scala:33 []
+-(4) MapPartitionsRDD[11] at map at PageRank.scala:59 []
| MapPartitionsRDD[10] at leftOuterJoin at PageRank.scala:59 []
| MapPartitionsRDD[9] at leftOuterJoin at PageRank.scala:59 []
| CoGroupedRDD[8] at leftOuterJoin at PageRank.scala:59 []
+-(4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []
| ShuffledRDD[7] at reduceByKey at PageRank.scala:54 []
+-(4) MapPartitionsRDD[6] at map at PageRank.scala:49 []
| MapPartitionsRDD[5] at values at PageRank.scala:49 []
| MapPartitionsRDD[4] at join at PageRank.scala:48 []
| MapPartitionsRDD[3] at join at PageRank.scala:48 []
| CoGroupedRDD[2] at join at PageRank.scala:48 []
| +-(4) ParallelCollectionRDD[0] at parallelize at PageRank.scala:33 []
| +-(4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []
| ShuffledRDD[31] at reduceByKey at PageRank.scala:54 []

```

```

+- (4) MapPartitionsRDD[30] at map at PageRank.scala:49 []
| MapPartitionsRDD[29] at values at PageRank.scala:49 []
| MapPartitionsRDD[28] at join at PageRank.scala:48 []
| MapPartitionsRDD[27] at join at PageRank.scala:48 []
| CoGroupedRDD[26] at join at PageRank.scala:48 []
+- (4) ParallelCollectionRDD[0] at parallelize at PageRank.scala:33 []
+- (4) MapPartitionsRDD[23] at map at PageRank.scala:59 []
| MapPartitionsRDD[22] at leftOuterJoin at PageRank.scala:59 []
| MapPartitionsRDD[21] at leftOuterJoin at PageRank.scala:59 []
| CoGroupedRDD[20] at leftOuterJoin at PageRank.scala:59 []
+- (4) MapPartitionsRDD[11] at map at PageRank.scala:59 []
| | MapPartitionsRDD[10] at leftOuterJoin at PageRank.scala:59 []
| | MapPartitionsRDD[9] at leftOuterJoin at PageRank.scala:59 []
| | CoGroupedRDD[8] at leftOuterJoin at PageRank.scala:59 []
| +- (4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []
| | ShuffledRDD[7] at reduceByKey at PageRank.scala:54 []
| +- (4) MapPartitionsRDD[6] at map at PageRank.scala:49 []
| | MapPartitionsRDD[5] at values at PageRank.scala:49 []
| | MapPartitionsRDD[4] at join at PageRank.scala:48 []
| | MapPartitionsRDD[3] at join at PageRank.scala:48 []
| | CoGroupedRDD[2] at join at PageRank.scala:48 []
| +- (4) ParallelCollectionRDD[0] at parallelize at PageRank.scala:33 []
| +- (4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []
| ShuffledRDD[19] at reduceByKey at PageRank.scala:54 []
+- (4) MapPartitionsRDD[18] at map at PageRank.scala:49 []
| MapPartitionsRDD[17] at values at PageRank.scala:49 []
| MapPartitionsRDD[16] at join at PageRank.scala:48 []
| MapPartitionsRDD[15] at join at PageRank.scala:48 []
| CoGroupedRDD[14] at join at PageRank.scala:48 []
+- (4) ParallelCollectionRDD[0] at parallelize at PageRank.scala:33 []
+- (4) MapPartitionsRDD[11] at map at PageRank.scala:59 []
| MapPartitionsRDD[10] at leftOuterJoin at PageRank.scala:59 []
| MapPartitionsRDD[9] at leftOuterJoin at PageRank.scala:59 []
| CoGroupedRDD[8] at leftOuterJoin at PageRank.scala:59 []
+- (4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []
| ShuffledRDD[7] at reduceByKey at PageRank.scala:54 []
+- (4) MapPartitionsRDD[6] at map at PageRank.scala:49 []
| MapPartitionsRDD[5] at values at PageRank.scala:49 []
| MapPartitionsRDD[4] at join at PageRank.scala:48 []
| MapPartitionsRDD[3] at join at PageRank.scala:48 []
| CoGroupedRDD[2] at join at PageRank.scala:48 []
+- (4) ParallelCollectionRDD[0] at parallelize at PageRank.scala:33 []
+- (4) ParallelCollectionRDD[1] at parallelize at PageRank.scala:42 []

```

What was executed by a job triggered by the program

The program makes use of `toDebugString` to see the lineage of the “ranks” RDD. Once transformation `leftOuterJoin` is used on ranks RDD, that’s when it is required to compute the RDD as there are actions to be performed on the RDD. This job triggers the execution of the ranks RDD.

Observations

1. Is Spark smart enough to figure out that it can re-use RDDs computed for an earlier action?

No. RDD just contains a reference to the task that has to be performed in the lineage. RDD does not store any data.

RDD operations are lazy. When an action is called, then through the lineage, each required RDD is computed and then used. Until action is called, RDDs are just tracked in lineage and are not computed. These RDDs are not stored in memory if not told by `cache()` or `persist()`. If `cache()` or `persist()` are not used, RDDs are not stored. For every action or iteration that requires a particular RDD, it's always computed again and again by executing the lineage. Since RDDs don't store data in them, re-use of RDDs are not possible.

2. How do `persist()` and `cache()` change this behavior?

`persist()` and `cache()` are methods in Spark to save the computed RDDs in memory.

These above two methods are lazy as well. Only when an action is used on RDD, then its computed and then if `cache()` or `persist()` are used, it's stored in memory.

By storing in memory, if the same computed RDD is required for any other action or in the next iteration, it can be fetched from memory and can avoid re-computation. Re-use of RDD is possible in this case. If memory is full with a lot of data from RDD, then spark automatically switches to re-computation of RDD, rather than storing it.

Difference between `cache()` and `persist()` is that `cache()` can store in `MEMORY_ONLY`. Whereas, with `persist()`, various storage levels can be chosen from.

When there is RDD partitioning required or re-use of RDD in iterations, it's better to use `cache()` or `persist()` to store it in memory.

Page Rank in Spark

Pseudo-code

To generate a graph (Map Only) –

```
Mapper {
    map(k)
    {
        for( int i=1 to k*k)                // create a graph from 1 to k^2
        {
            if(i%k==0)                      // for every kth vertex outgoing link doesn't exists
                emit(i , 0)
            else
                emit(i, i+1)                // all other vertices have outgoing link to next vertex
        }
    }
}
```

To generate adjacency list with initial page rank

Input in form – vertex1, vertex2

vertex1, vertex3

Output in form – vertex1 vertex2#vertex3;page_rank

```
Mapper {
    map()
    {
        String s = covert value to string
        String[] followers = s.split(",")    // split string on delimited as ","
    }
}
```

```

    write a key-value pair as (followers[0],followers[1]) to context
}
}
Reducer {
reduce(key,values)
{
    String adj = "";
    for each value of values for the key
    {
        adj += node + "#"; // every adjacent vertex is separated by #
    }
    String outputVal = adj+";"+(1/k*k) // value is adjacency list and initial pagerank separated by ";"
    Context.write(key,outputVal)
}
}

```

To Calculate Page Rank for n iterations(2 jobs)

1st Job to calculate page rank

2nd Job to add dangling page mass to each vertex's PR

```

Counter PR_Mass{
    dangling_mass //counter to keep track of dangling mass
}

```

JOB 1

```

Mapper {
    map((id n, vertex N)
    {
        emit(n, N) // Pass along the graph structure
        p = N.pageRank / N.adjacencyList.size() // Compute contributions to send along outgoing links
        for all m in N.adjacencyList do
            emit( m, p )
        }
    }
}

Reducer {
    reduce((id m, [p1, p2,...])
    {
        If(m!=0)
        {
            s = 0
            M = NULL
            for all p in [p1, p2,...] do
                if isVertex(p) then // The vertex object was found: recover graph structure
                    M = p
                Else
                    s += p // A PageRank contribution from an inlink was found:
                        // add it to the running sum.
            M.pageRank = (1-α)/|V| + α*s // PR formula
            emit( m, M)
        }
    }
}

```

```

        Else
        {
            s =0
            for all p in [p1, p2,...] do
                s+=p
            PR_mass.dangling_mass = s / |V|
        }
    }
}

```

// if it's vertex 0 then add the dangling mass and
// update counter

JOB 2

```

Mapper {
    map((id n, vertex N)
    {
        emit(n, N)
    }
}

Reducer {
    dMass = PR_mass.dangling_mass
    reduce((id m, [p]))
    {
        String[ ] adj = p.split(";")
        Float newPR = adj[1].toInt()+dMass
        String outputVal = adj[0]+";"+newPR
        emit(m,outputVal)
    }
}

```

// Pass along the graph structure

// dangling mass to be added to each vertex's PR

// split value by ";" to get the PR
// add dMass to old PR

To add Page Ranks of every vertex

```

Mapper
{
    map((id n, vertex N)
    {
        String[] inputStr = s.split("\t")
        String[] PRStr = inputStr.split(";")
        String emitPR = PRStr[1]
        emit("PR", emitPR)
    }
}

Reducer
{
    reduce("PR",[p1,p2,p3....])
    {
        finalPR =0
        For p in all values do
            finalPR += p
        emit("final PR", finalPR)
    }
}

```

// split string on delimited as "\t"
// split adjacency list and PR
// consider page rank only
// emit same key for all and value as PR

How was dangling-page problem handled?

Dummy vertex is add with key 0 for every vertex which doesn't have an outgoing link. In the reducer job, all page rank values of 0 vertices is aggregated and divided by the number of vertices in the graph. That result called the dangling mass is updated to every vertex's page rank value.

As shown above in the pseudo code, this is done with 2 jobs.

1 job to only calculate the page rank for every vertex ignoring the dangling mass.

2nd job to add the dangling mass to every vertex.

This dangling mass is updated in the global counter in the reducer of the 1st job. Dangling page mass is updated to the configuration object in the driver program. And the updated dangling mass is used in reducer of the 2nd job where, for each vertex, (dangling mass * α) is added to it's page rank.

This way the dangling mass is distributed to each vertex.

Usage of NLineInputFormat

Used NLineInputFormat to pass 50000 lines as input for each Mapper.

As there will be 1000,000 line in the file which holds the Adjacency list and page rank, each Mapper receives 50000 lines of input.

With this, for every job which uses NLineInputFormat class:

Number of Map tasks = 22

Number of Reduce tasks = 12

NLineInputFormat class is used in the following files of the code: PRCalc.java, PageRank.java, PageRankAgg.java .

Running Time Measurements

Map-Reduce:

For k = 1000

	<i>6 machines</i>	<i>11 machines</i>
<i>Start Time</i>	00:49:05	01:40: 38
<i>End Time</i>	01:13:42	01:59:11
<i>Time Taken (in sec)</i>	24 min 37 sec	18 min 33 sec
<i>Page Rank of 1000 pages</i>	0.9946458287277518	0.9946458287277516